



Figure 1: Comparison between the R-PAGE and R-SPIDER methods (left plot), and between R-MARINA and Riemannian Compressed SGD (right plot). For the compressed methods (right plot), we employ the Rand- $K$  compressor, retaining 30% of the vector entries.